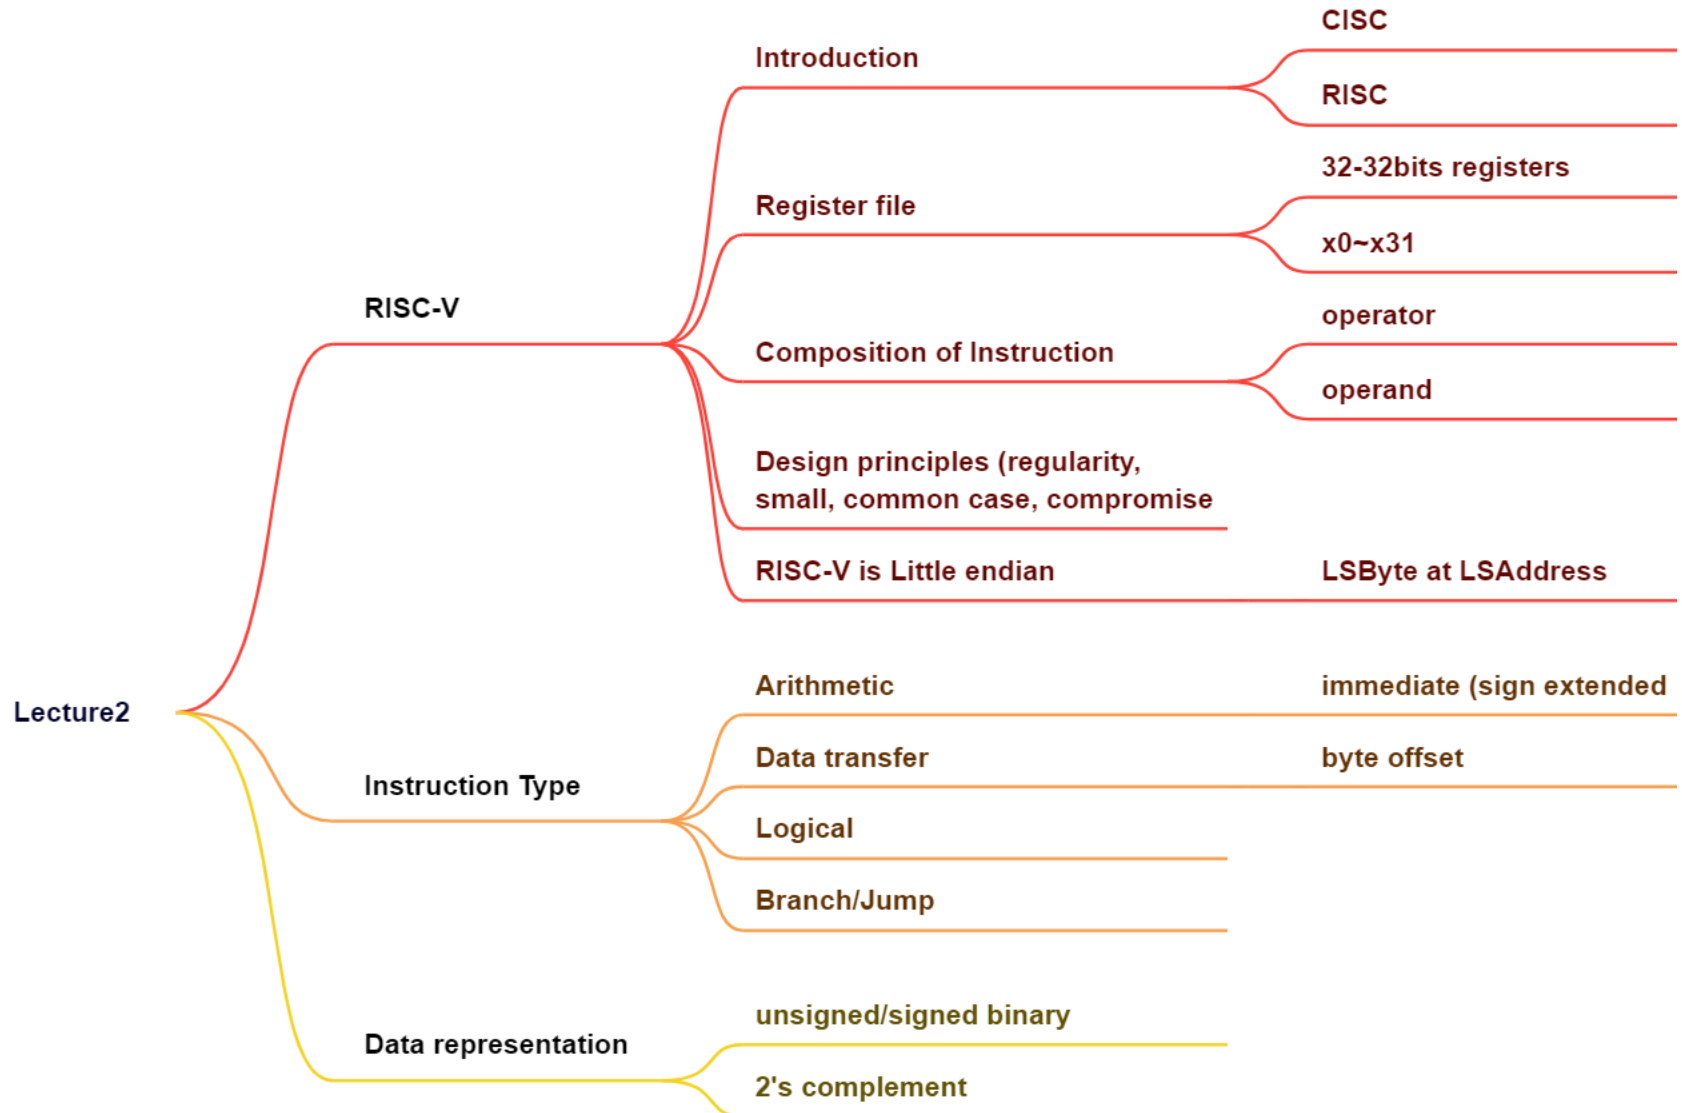


# COMPUTER ORGANIZATION

## Lecture 3 RISC-V Procedure

2024 Spring

# Recap



# Recap: Exercise

- For the following C statement, what is the corresponding RISC-V RV32 assembly code? Assume that the variable `i` is assigned to register `x13`, respectively. Assume that the base address of the arrays `A` and `B` are in registers `x16` and `x17`, respectively.

`B[8] = A[i-3];`

# Conditional Operations

- Branch to a labeled instruction if a condition is true
  - Otherwise, continue sequentially
- Conditional branch
  - *beq rs1, rs2, L1*
    - if (rs1 == rs2) branch to instruction labeled L1;
  - *bne rs1, rs2, L1*
    - if (rs1 != rs2) branch to instruction labeled L1;
- Unconditional branch
  - *beq x0, x0, L1*
    - unconditional jump to instruction labeled L1

# Labels in Assembly


- We commonly see "labels" in the code
  - `foo: add x2, x1, x0`
- The assembler converts these into positions in the code
  - At what address in the code is that label ...
- Labels give control flow instructions, such as jumps and branches, a place to go ...
  - e.g. `bne x0, x2, foo`
- The assembler in outputting the code does the necessary calculation so the jump or branch will go to the right place

# Compiling If Statements

- C code

```
if (i==j) f = g+h;
else f = g-h;
```

- i and j are in x22 and x23,
- f,g and h are in x19, x20 and x21



- Compiled RISC-V code:

```

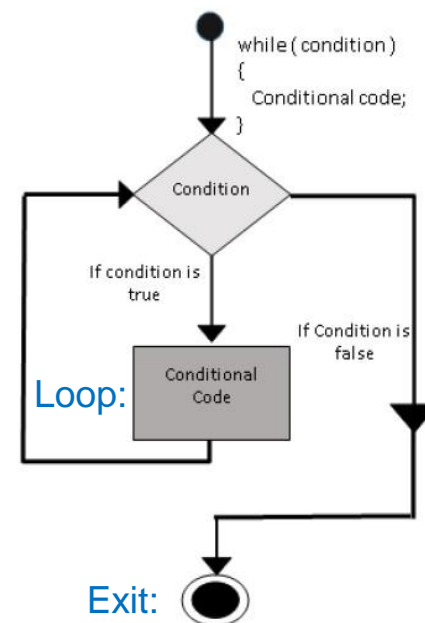

    bne x22, x23, Else # go to Else if i ≠ j
    add x19, x20, x21 # f=g+h, skipped if i ≠ j
    beq x0, x0, Exit  # unconditional go to Exit
Else: sub x19, x20, x21 # f=g-h, skipped if i = j
Exit:
```

# Compiling Loop Statements

- C code:

```
while (save[i] == k)
    i += 1;
```

- i in x22, k in x24, address of save in x25



- Compiled MIPS code:

```
Loop: sll    x10, x22, 2      # Temp reg x10 = i * 4
      add    x10, x10, x25    # x10 = address of save[i]
      lw     x9, 0(x10)       # Temp reg x9 = save[i]
      bne    x9, x24, Exit    # go to Exit if save[i]≠k
      addi   x22, x22, 1      # i = i + 1
      j      Loop            # go to Loop
```

Exit:

# More Conditional Operations

- Signed comparison
  - `blt rs1, rs2, L1`
    - if ( $rs1 < rs2$ ) branch to instruction labeled L1
  - `bge rs1, rs2, L1`
    - if ( $rs1 \geq rs2$ ) branch to instruction labeled L1
  - Example, C to RISC-V
    - `if (a > b) a += 1;`
    - a in x22, b in x23
      - `bge x23, x22, Exit # signed comparison`
      - `addi x22, x22, 1`
  - Exit:
- Unsigned comparison
  - `bltu, bgeu`



# What if we need more instructions?

- RISC-V doesn't have "branch if greater than" or "branch if less than or equal"
- Instead you can reverse the arguments, as:
  - $A > B$  is equivalent to  $B < A$
  - $A \leq B$  is equivalent to  $B \geq A$
- The assembler defines **pseudo-instructions** for your convenience:  
`bgt x2, x3, foo (pseudo)` will become  
`blt x3, x2, foo (basic)`


# Pseudo-instructions

- For more pseudo-instructions, refer to RARS Help (see in lab).

Basic Instructions	Extended (pseudo) Instructions	Directives	Syscalls	Exc
lhu t1,10000000	Load Halfword Unsigned : Set t1 to zero-extended 16-bit value			
lhu t1,label	Load Halfword Unsigned : Set t1 to zero-extended 16-bit value			
li t1,-100	Load Immediate : Set t1 to 12-bit immediate (sign-extended)			
li t1,10000000	Load Immediate : Set t1 to 32-bit immediate			
lui t1,%hi(label)	Load Upper Address : Set t1 to upper 20-bit label's address			
lw t1,%lo(label) (t2)	Load from Address			
lw t1,(t2)	Load Word : Set t1 to contents of effective memory word address			
lw t1,-100	Load Word : Set t1 to contents of effective memory word address			
lw t1,10000000	Load Word : Set t1 to contents of effective memory word address			
lw t1,label	Load Word : Set t1 to contents of memory word at label's address			
mv t1,t2	MoVe : Set t1 to contents of t2			
neg t1,t2	NEGate : Set t1 to negation of t2			
nop	NO OPeration			
not t1,t2	Bitwise NOT (bit inversion)			

# Basic Blocks

- A basic block is a sequence of instructions with
  - No embedded branches (except at end)
  - No branch targets (except at beginning)



- A compiler identifies basic blocks for optimization
- An advanced processor can accelerate execution of basic blocks

# C to RISC-V Example

```
int A[20];
int sum = 0;
for (int i=0; i<20; i++)
    sum += A[i];
```

Loop has 7 instructions

```
# Assume x8 holds pointer to A
# Assign x10=sum, x11=i
1  add x10, x0, x0          # sum=0
2  add x11, x0, x0          # i=0
3  addi x12,x0,20           # x12=20
4  loop:
   bge x11, x12, exit
5  slli x13, x11, 2         # i * 4
6  add x13, x13, x8         # A + i
7  lw x13, 0(x13)          # *(A + i)
8  add x10, x10, x13        # increment sum
9  addi x11, x11, 1         # i++
10 beq x0, x0, loop        # iterate
11 exit:
```

# C to RISC-V Example Optimized

```
int A[20];  
int sum = 0;  
for (int i=0; i<20; i++)  
    sum += A[i];
```

Loop now has 6 instructions

```
# Assume x8 holds base address of A  
# Assign x10=sum, x11=i*4  
1  add x10, x0, x0  # sum=0  
2  add x11, x0, x0  # i=0  
3  addi x12,x0,80   # x12=20*4  
4  loop:  
   bge x11, x12, exit  
5  add x13, x11, x8  # A + i  
6  lw x13, 0(x13)    # *(A + i)  
7  add x10, x10, x13  # increment sum  
8  addi x11, x11, 4   # i++  
9  beq x0, x0, loop  # iterate  
10 exit:
```

# C to RISC-V Example Optimum

```
int A[20];  
int sum = 0;  
for (int i=0; i<20; i++)  
    sum += A[i];
```

Loop now has 4 instructions

- Directly increment ptr into A array
- And only 1 branch/jump rather than two
  - Because first time through is always true so can move check to the end
  - The compiler will often do this automatically for optimization

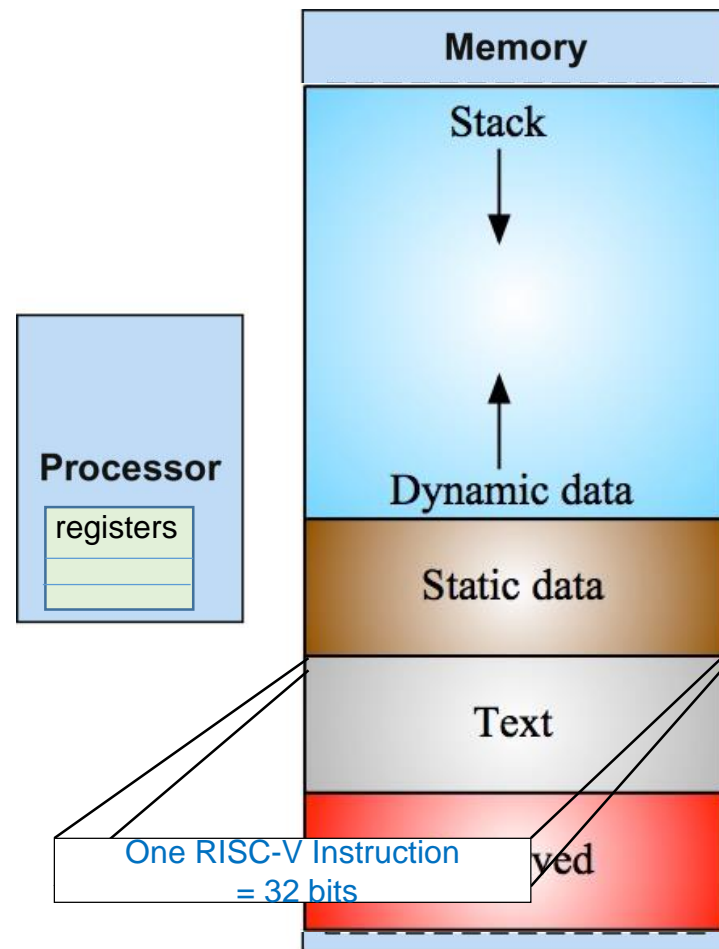
```
# Assume x8 holds base address of A  
# Assign x10=sum  
# Assume x11 holds ptr to next A  
1 add x10, x0, x0          # sum=0  
2 add x11, x0, x8          # Copy of A  
3 addi x12, x8, 80         # x12=80 + A  
4 loop:  
  lw x13, 0(x11)  
5 add x10, x10, x13  
6 addi x11, x11, 4  
7 blt x11, x12, loop
```

# Procedure Calling

- A procedure or function is one tool used by the programmers to structure programs
  - Benefit: easy to understand, reuse code
- We can think of a procedure like a spy
  - acquires resources → performs task → covers his tracks → returns back with desired result
- Six Steps of Calling a Function
  1. Put parameters **in a place** where the procedure can access them.
  2. **Transfer control** to the procedure.
  3. Acquire the **storage resources** needed for the procedure.
  4. **Perform** the desired task.
  5. Put the result value **in a place** where the calling program can access it.
  6. **Return control** to the **point of origin**, since a procedure can be called from several points in a program.

# Recall: How Program is Stored

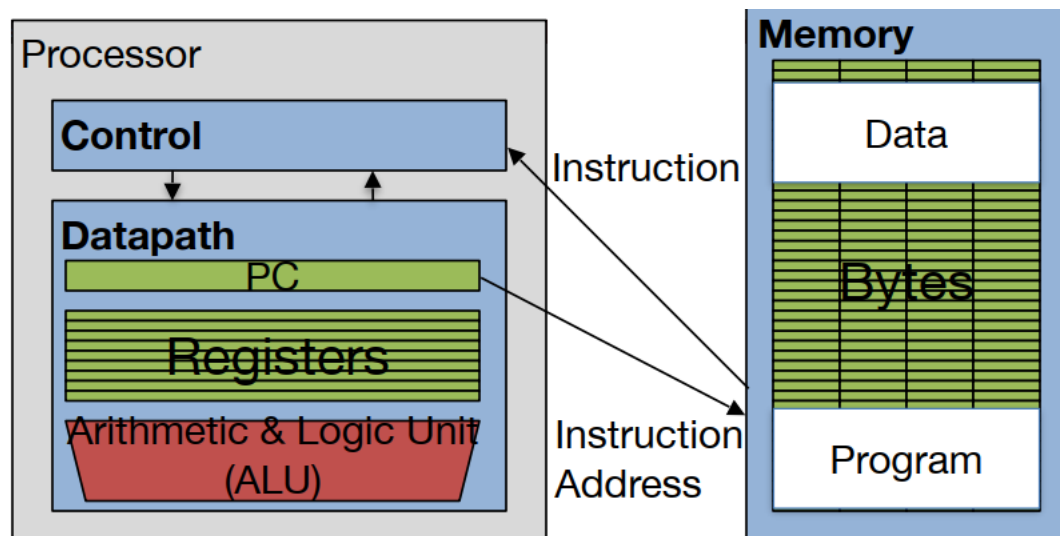
- Instructions(programs) are represented in binary, just like data
- Programs are stored in *Text Segment*
- Constants and other static variables are stored in *Static data segment*
- Dynamic data: *Heap*
  - E.g., malloc in C, new in Java
- Automatic data: *Stack*





# Program Execution

- **PC (program counter)** is special internal register inside processor holding **byte** address of next instruction to be executed
- Instruction is fetched from memory, then control unit executes instruction using datapath and memory system, and updates program counter (default is add +4 bytes to PC, to move to next sequential instruction)



# Procedure Calling

Caller:

```
int x = 5;
int y = 3;
int z = sum (x , y);
x = x + 7;
...
```

callee:

```
int sum ( int a, int b)
{
    int c = a + b;
    return c;
}
```

Data Memory

s0(x8)
s1(x9)
s2(x18)
a0(x10)
a1(x11)
a2(x12)
ra(x1)

Register

⋮
<del>5</del> 8
3
⋮
⋮
⋮


a0(x10)

a1(x11)

PC

Program Memory

x = 5;
y = 3;
int z = sum (x , y);
...
int sum ( int a, int b)
{ int c = a + b;
return c; }



# Procedure Calling: Question1

- Step 1, 5 and 6: Where should we put the arguments and return values?
  - Registers way faster than memory, so use them whenever possible
  - Symbolic register names
    - E.g., `a0-a7 (x10-x17)` for function arguments, `a0-a1` for return values
    - E.g., `ra (x1)` for return address, used to save where a function is called from so we can get back
  - If need extra space, use memory (the stack!)

# RISC-V Registers and Convention

- ABI: Application Binary Interface, defines “Calling Convention” – How to call other functions
- So going forward, no more x3, x6... type nomenclature

Register	ABI name	Description	Saved by Callee?
x0	zero	the constant value 0	N/A
x1	ra	Return address	No
x2	sp	Stack pointer	Yes
x3	gp	Global pointer	N/A
x4	tp	Thread pointer	N/A
x5 – 7	t0 – 2	Temporaries	No
x8	s0 / fp	Saved register/Frame pointer	Yes
x9	s1	Saved register	Yes
x10 – 17	a0 – 7	Function arguments/Return values	No
x18 – 27	s2 – 11	Saved registers	Yes
x28 – 31	t3 – 6	Temporaries	No

# Procedure Calling: Question2

- Step 2 and 6: How do we Transfer Control?
  - Procedure call: **jump and link** (the real order is link and jump)  
`jal ra, target` #pseudo code: `jal target`
    - Address of following instruction (return address) put in ra
    - Jumps to target address
    - Used by Caller
  - Procedure return: **jump register**  
`jalr zero, 0(ra)` #pseudo code: `jr ra`
    - Similar to “Jump and Link” except in specification of target
    - Jump to ra and simultaneously saves the address of following instruction in zero register (value put into zero is meant to be throw away)
    - Used by Callee

# Transfer Control

Caller:

```
int x = 5; // x in s0
int y = 3; // y in s1
int z = sum (x , y);
...
```

callee: // a in a0, b in a1, c in a0

```
int sum ( int a, int b) {
    int c = a + b;
    return c;
}
```

address (shown in decimal)

```
1000 mv a0,s0      # parse argument x → a
1004 mv a1,s1      # parse argument y → b
1008 jal sum       # ra=1012, goto sum
1012 ...           # next instruction
...
```

```
2000 sum: add a0,a0,a1
```

```
2004 jr ra        # next instr at 1012...
```

Note: these are pseudo code

PC's decimal value: ...→1000→1004→1008→2000→2004→1012→...

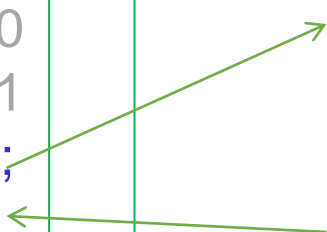
# Transfer Control

Caller:

```
int x = 5; // x in s0
int y = 3; // y in s1
int z = sum (x , y);
...
```

callee: // a in a0, b in a1, c in a0

```
int sum ( int a, int b) {
    int c = a + b;
    return c;
}
```



- Question: Why use **jr** here? Why not use **j**?
- Answer: **sum** might be called by many places, so we can't return to a fixed place. The calling proc to **sum** must be able to say "return here" somehow.

2000 **sum:** add a0, a0, a1

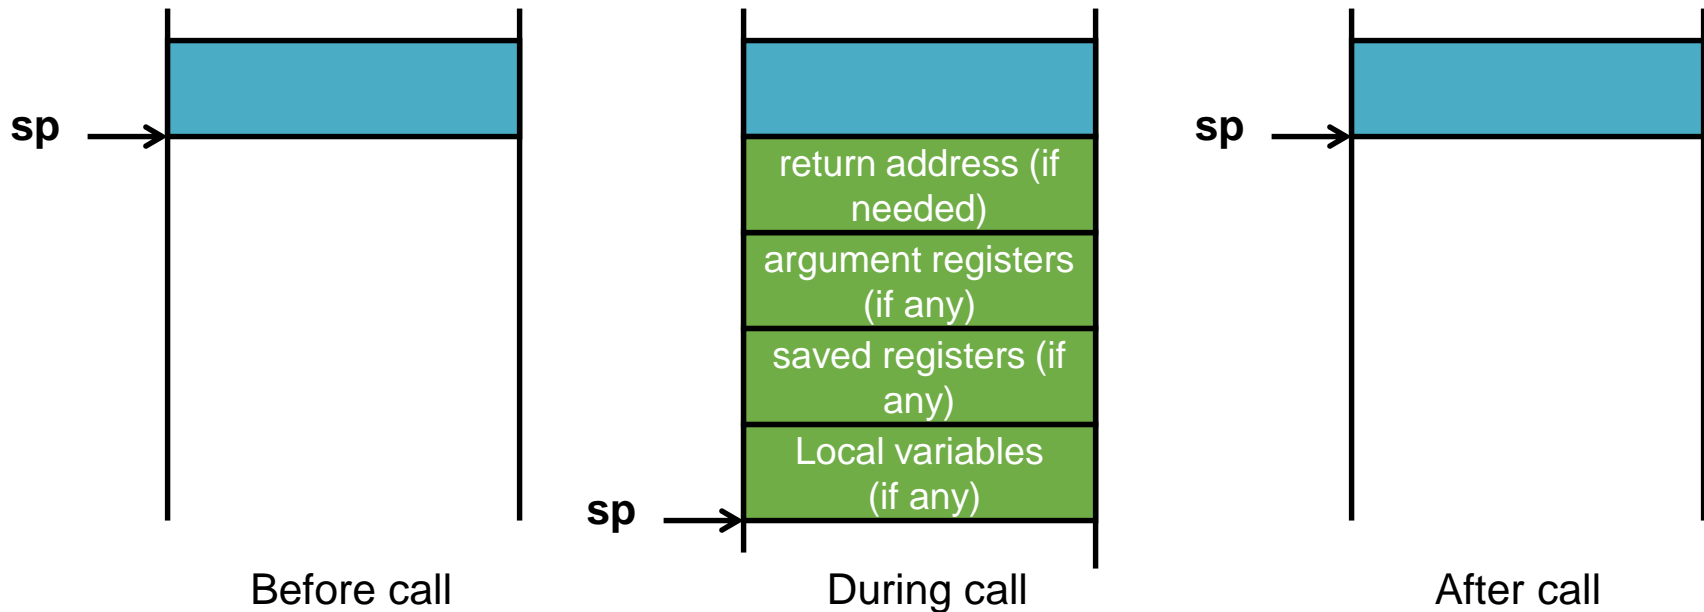
2004 **jr** ra # jump to the return address

# Procedure Calling: Question3

- Step 3: What's the Local storage for variables?
  - C has three storage classes:
    - Stack: automatic variables are local to function and discarded when function exits
    - Static: variables exist across exits from and entries to procedures
    - Heap: variables declared dynamically via malloc
  - Use stack for automatic (local) variables that don't fit in registers
    - Push: placing data onto stack
    - Pop: removing data from stack
  - **sp (x2) (stack pointer)** points to the last used place at the stack
    - Push decrements sp
    - Pop increments sp



# Stack Before, During, After Call



# Procedure Calling: Question 4

- Step 4: Function Calling Conventions?
  - It is effectively a contract between functions
  - By convention, registers are classified as one of ...
  - **caller-saved**
    - The function invoked (the callee) can do whatever it wants to them!
    - Means that the caller can not count on their contents not being destroyed
  - **callee-saved**
    - The function invoked must restore them before returning (if used)

# The Calling Convention

- Caller saved:
  - **a0–a7 (x10-x17)**: eight argument registers to pass parameters and two return values (**a0-a1**)
  - **t0-t6** Temporaries
  - **ra**: one return address register for return to the point of origin
- Callee saved:
  - **s0-s11** Saved registers: Preserved across function calls
- If both the Caller and Callee obey the procedure conventions, there are significant benefits
  - People who have never seen or even communicated with each other can write functions that work together
  - Recursion functions work correctly

# Leaf Procedure Example

- A “leaf” function - it calls nothing

```
int Leaf(int g, int h, int i, int j) {  
    int f;  
    f = (g + h) - (i + j);  
    return f;  
}
```

- Parameter variables g, h, i, and j in argument registers a0, a1, a2, and a3.
- Return variable f uses register a0
- Assume we compute f by using s0 and s1
- s0 and s1 are callee saved, so it's the responsibility of “Leaf” to save and restore

# Leaf Procedure Example

a0, a1, a2, a3

- RISC-V code:

```
int Leaf(int g, int h, int i, int j)
{ int f;
  f = (g + h) - (i + j);
  return f;
}
```

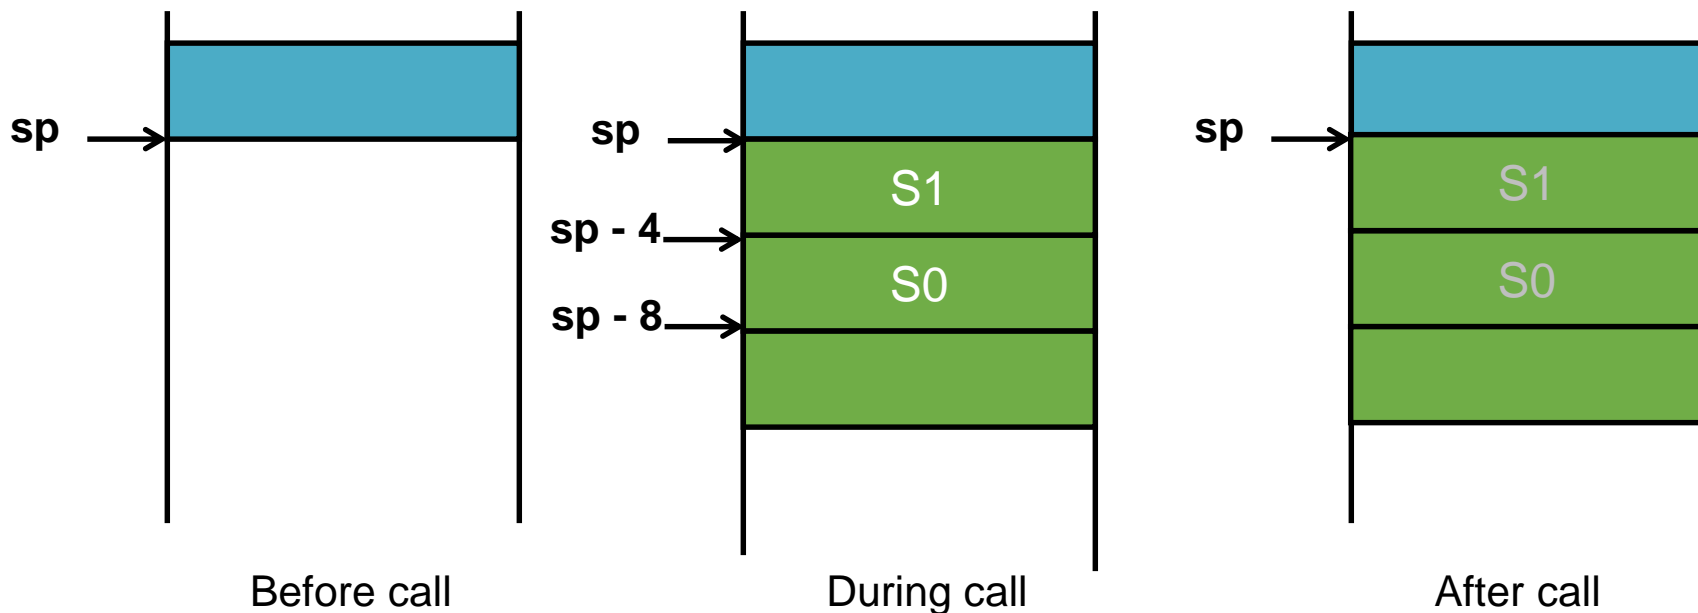
Diagram showing register mappings:   
 - **a0** points to the parameter **g**.   
 - **s0** points to the variable **f**.   
 - **s1** points to the expression **(i + j)**.

Leaf:

1	addi sp,sp,-8	# adjust stack for 2 items
2	sw s1, 4(sp)	# save s1 for use afterwards
3	sw s0, 0(sp)	# save s0 for use afterwards
4	add s0,a0,a1	# s0 = g + h
5	add s1,a2,a3	# s1 = i + j
6	sub a0,s0,s1	# return value (g + h) - (i + j)
7	lw s0, 0(sp)	# restore register s0 for caller
8	lw s1, 4(sp)	# restore register s1 for caller
9	addi sp,sp,8	# adjust stack to delete 2 items
10	jr ra	# jump back to calling routine # pseudo of jalr zero, 0(ra)

# Leaf Procedure Example: Stack

- In the previous example we need to save old values of `s0` and `s1`
- Push doesn't actually delete content from memory, it just increase the stack pointer



# Leaf Procedure Example: Observation

- This is a "leaf function": it calls no other function
  - We didn't need to save ra (because leaf didn't call any other function and therefore ra never changed)
  - Instead of s0 and s1 can just use temporary (caller-saved registers) only
- So we could have just as easily used t0 and t1 instead...

leaf:

```
add t0,a0,a1 # t0 = g + h
add t1,a2,a3 # t1 = i + j
sub a0,t0,t1 # return value (g+h)-(i+j)
jr ra
```

Caller saved:  
a0–a7  
t0–t6  
ra  
Callee saved:  
s0–s11

# Non-Leaf Procedures

- Procedures that call other procedures (nested call)
  - Note: this could mean a function calling itself – recursion
- For nested call, caller needs to save on the stack:
  - Its return address
  - Any arguments and temporaries needed after the call
- Restore from the stack after the call



# Non-Leaf Procedure Example

- C code:

```
int sumSquare(int x, int y) {  
    return mult(x,x)+ y;  
}
```

- Function **sumSquare** is calling **mult**
- So there's a value in **ra** that **sumSquare** wants to jump back to, but this will be overwritten by the call to **mult**
- Need to save **sumSquare** return address before call to **mult**

# Non-Leaf Procedure Example

a0, a1

```
int sumSquare(int x, int y) {  
    return mult(x,x)+ y;  
}
```

a0, a1

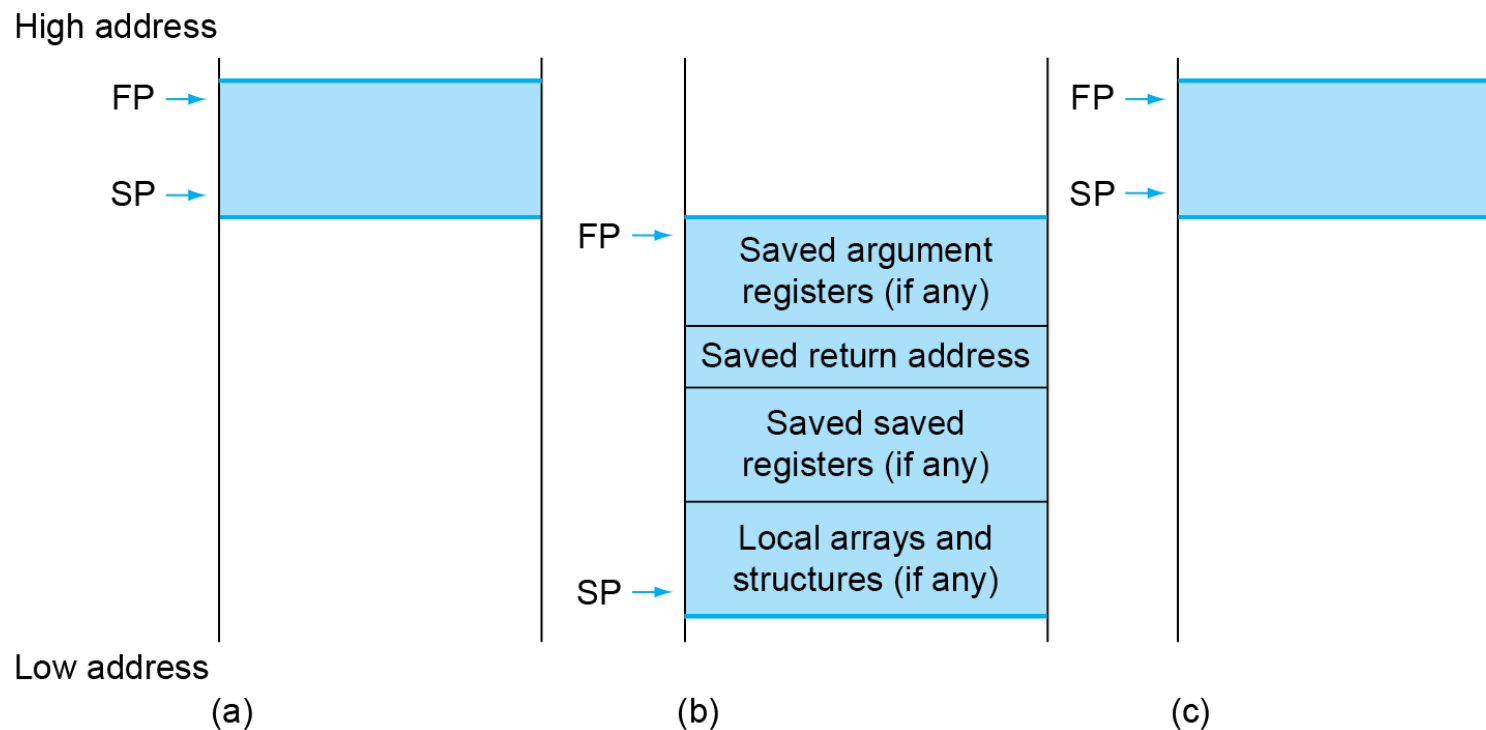
sumSquare:

```
1  addi sp,sp,-8 # reserve space on stack  
2  sw ra, 4(sp) # save ret addr Push  
3  sw a1, 0(sp) # save y  
4  mv a1,a0 # prepare 2nd argument for mult  
5  jal mult # call mult  
6  lw a1, 0(sp) # restore y  
7  add a0,a0,a1 # mult()+y  
8  lw ra, 4(sp) # get ret addr  
9  addi sp,sp,8 # restore stack Pop  
10 jr ra  
    mult:
```

# Register Conventions Summary

- Caller must save arguments and temporary registers it is using onto the stack before making a procedure call
- Caller can trust saved registers to maintain values
- Callee must save any saved registers it intends to use by putting them on the stack before overwriting their values
- Notes:
  - Caller and callee only need to save the **appropriate using** (not all!)
  - Don't forget to restore the values later

# Local Data on the Stack



- Local data allocated by callee
  - e.g., C automatic variables
- Procedure frame (activation record)
  - Used by some compilers to manage stack storage

# Summary

- Functions called with `jal`, return with `jalr (jr)`
- The stack is your friend: Use it to save anything you need. Just leave it the way you found it!
- Instructions we know so far...
  - Arithmetic
  - Memory
  - Conditional Branches
  - Unconditional Branches (Jumps)
- Registers

# Non-Leaf Procedure Example 2

```
int fact (int n)
{
    if (n < 1) return (1);
    else return (n * fact(n-1));
}
```

## Notes:

The caller saves a0 and ra in its stack space.

Temps are never saved.

Compare  $n < 1$

Return 1

Fact(n-1)

Return  $n * \text{fact}(n-1)$

```
fact:
    addi    sp, sp, -8
    sw      ra, 4(sp)
    sw      a0, 0(sp)
    addi    t0, a0, -1
    bge     t0, zero, L1
    addi    a0, zero, 1
    addi    sp, sp, 8
    jr      ra
```

L1:

```
    addi    a0, a0, -1
    jal     fact
    addi    t1, a0, 0
    lw      a0, 0(sp)
    lw      ra, 4(sp)
    addi    sp, sp, 8
    mul     a0, a0, t1
    jr      ra
```

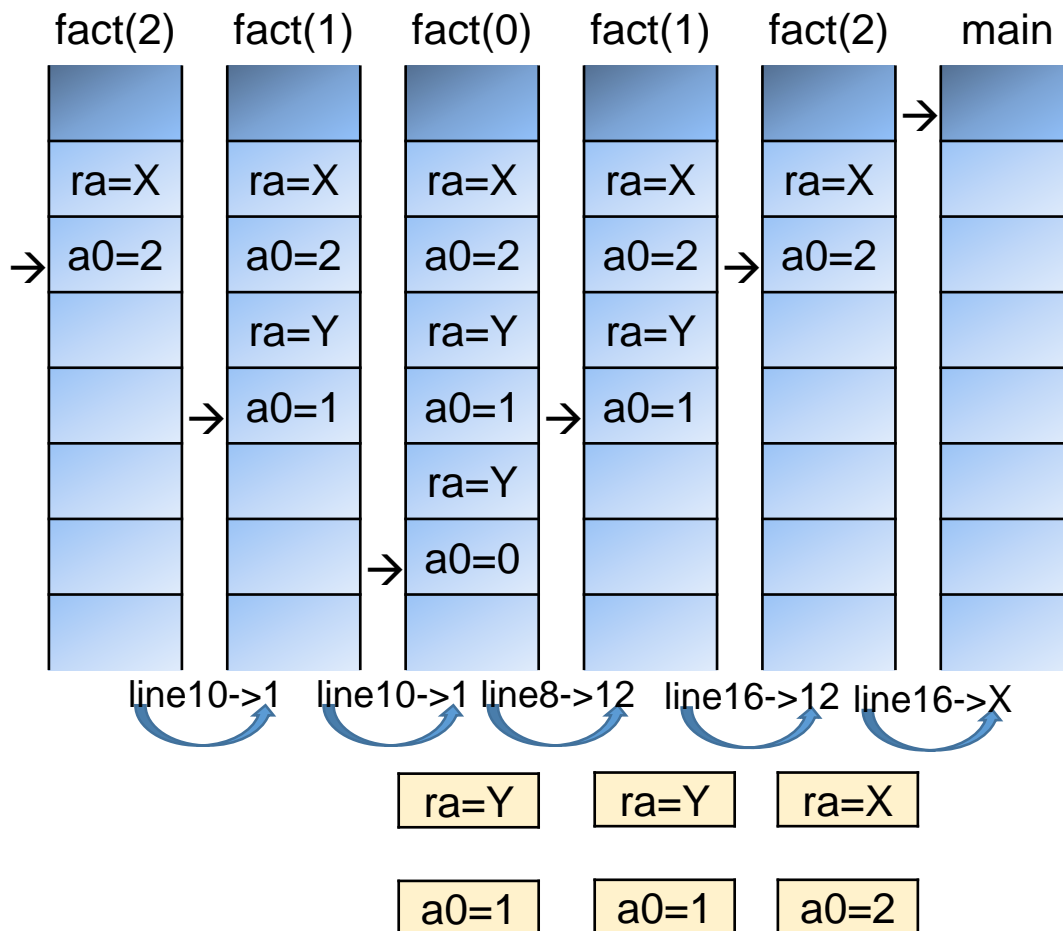
# Non-Leaf Procedure Example 2

- RISC-V code: suppose  $n=2$

```
int fact (int n)
{ if (n < 1) return 1;
  else return n * fact(n - 1); }
```

**fact:**

```
1  addi  sp, sp, -8
2  sw    ra, 4(sp)
3  sw    a0, 0(sp)
4  addi  t0, a0, -1
5  bge   t0, zero, L1
6  addi  a0, zero, 1
7  addi  sp, sp, 8
8  jr    ra
9  L1:   addi a0, a0, -1
10  jal   fact
11  mv    t1, a0
12 Y:    lw    a0, 0(sp)
13      lw    ra, 4(sp)
14      addi  sp, sp, 8
15      mul   a0, a0, t1
16      jr    ra
```



X: return address of caller  
(not shown in the code)

Y: the address of line 12