

COMPUTER ORGANIZATION

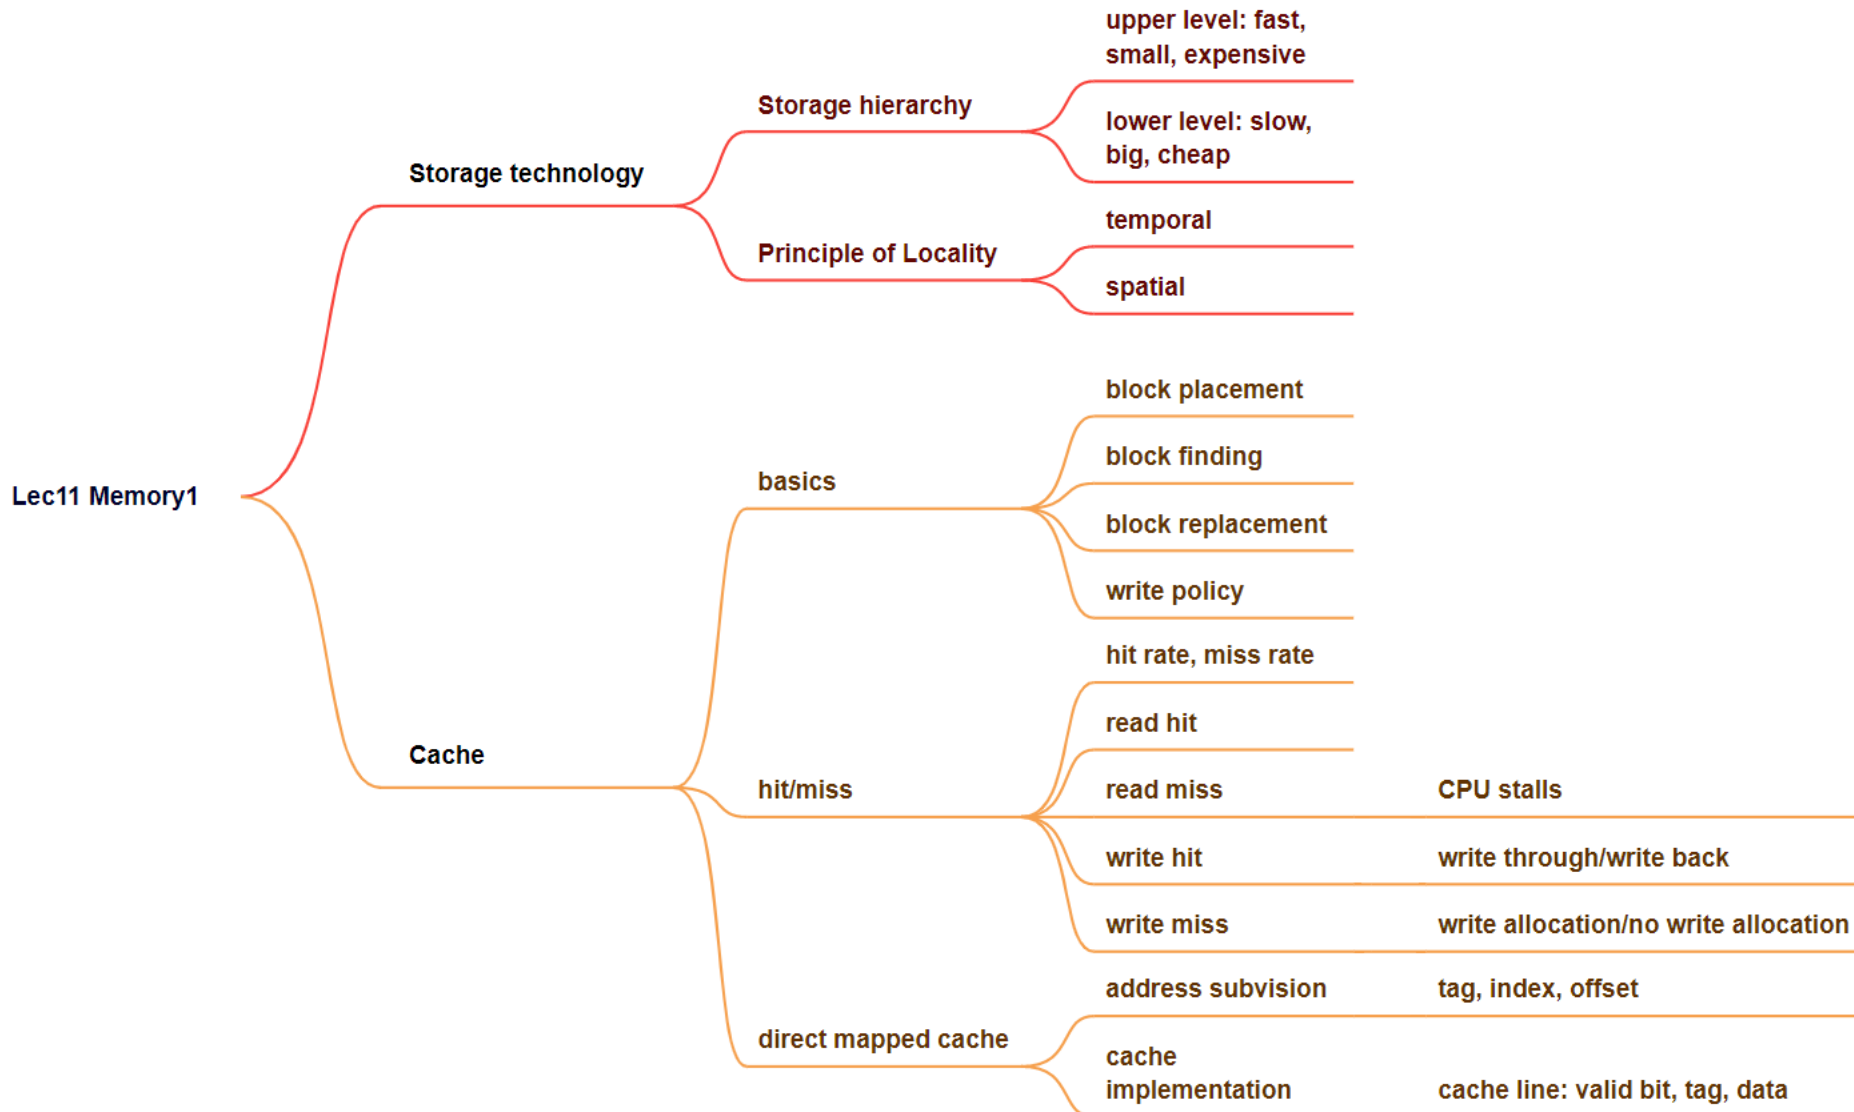
Lecture 12 Memory Hierarchy (2)

2024 Spring

This PowerPoint is for internal use only at Southern University of Science and Technology.
Please do not repost it on other platforms without permission from the instructor.

For Internal Use Only!

Recap



Outline

- **Measuring cache performance**
- Improving performance – Associative cache
 - Fully associative
 - n-ways Set associative
- Improving performance – Multilevel Caches



Measuring Cache Performance

- Components of CPU time
 - Program execution cycles
 - Includes cache hit time
 - Memory stall cycles
 - Mainly from cache misses
- With simplifying assumptions:

$$\begin{aligned} & \text{Memory stall cycles} \\ &= \frac{\text{Memory accesses}}{\text{Program}} \times \text{Miss rate} \times \text{Miss penalty} \\ &= \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Misses}}{\text{Instruction}} \times \text{Miss penalty} \end{aligned}$$



Cache Performance Example

- Given
 - I-cache miss rate = 2%
 - D-cache miss rate = 4%
 - Miss penalty = 100 cycles
 - Base CPI (ideal cache) = 2
 - Load & stores are 36% of instructions
- Miss cycles per instruction
 - I-cache: $2\% \times 100 = 2$
 - D-cache: $36\% \times 4\% \times 100 = 1.44$
- Effective CPI = $2 + 2 + 1.44 = 5.44$
 - Ideal CPU is $5.44/2 = 2.72$ times faster



Average Access Time

- Hit time is also important for performance
- Average memory access time (**AMAT**)

$$\text{AMAT} = \text{Hit time} + \text{Miss rate} \times \text{Miss penalty}$$

- Example
 - CPU with 1ns clock, hit time = 1 cycle, miss penalty = 20 cycles, I-cache miss rate = 5%
 - $\text{AMAT} = 1 + 0.05 \times 20 = 2\text{cycles} = 2\text{ns}$
 - 2 cycles per instruction

Performance Summary

- When CPU performance increased
 - Miss penalty becomes more significant
- Decreasing base CPI
 - Greater proportion of time spent on memory stalls
- Increasing clock rate
 - Memory stalls account for more CPU cycles
- Can't neglect cache behavior when evaluating system performance

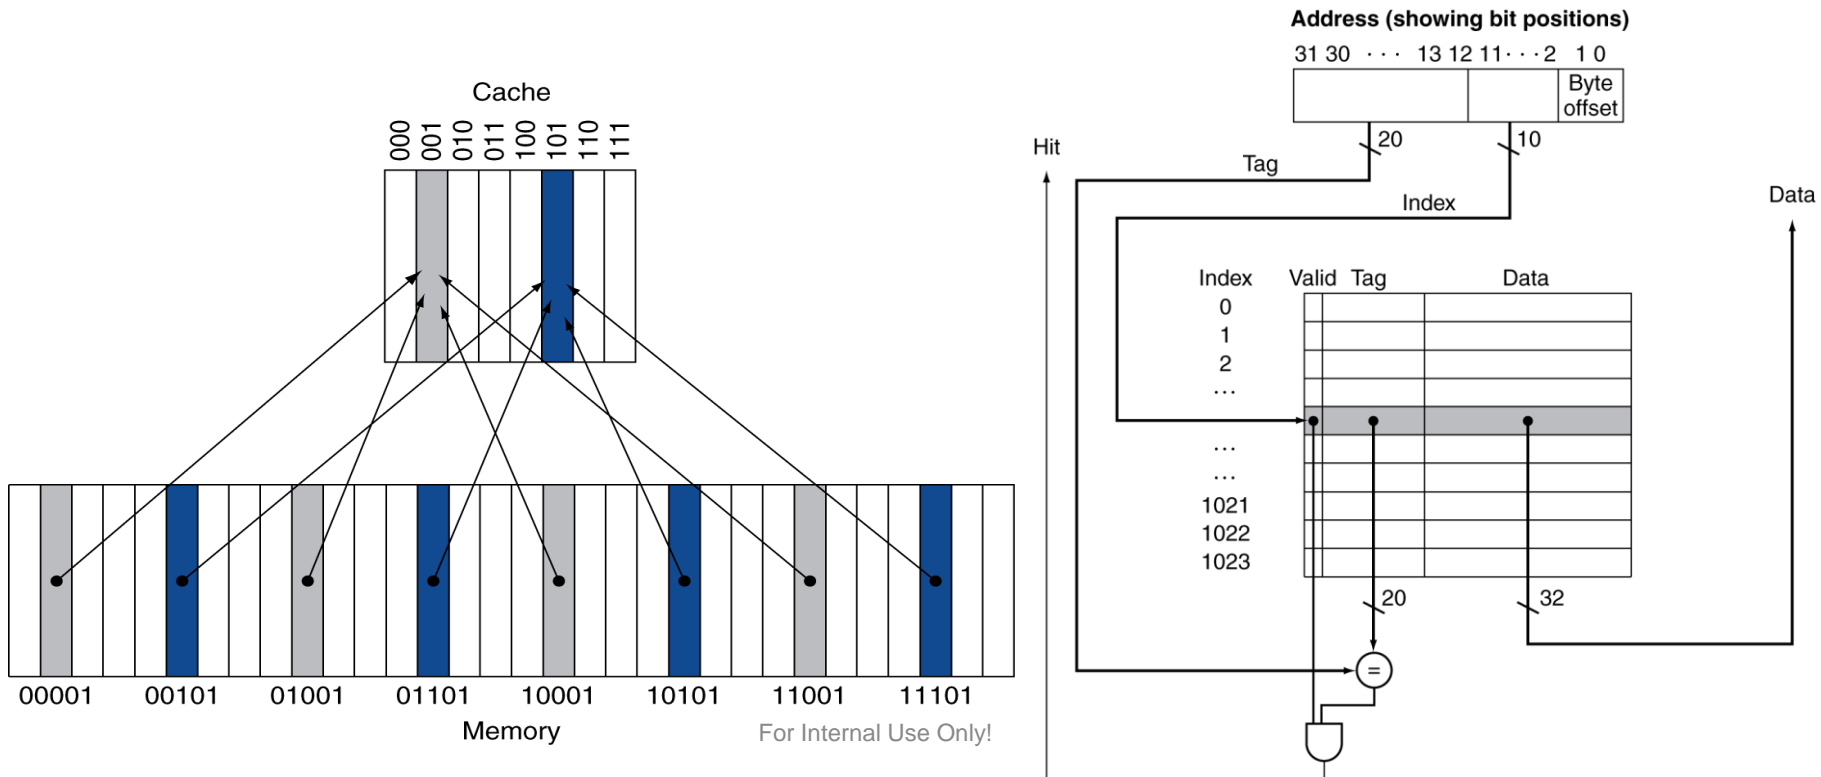


Outline

- Measuring cache performance
- **Improving performance – Associative cache**
 - Fully associative
 - n-ways Set associative
- Improving performance – Multilevel Caches

Recall: Direct Mapped Cache

- Direct mapped cache:
 - Location determined by address
 - One data in memory is mapped to **only one location** in cache
 - Capacity of cache is not fully exploited
 - Miss rate is high



Recall: Direct Mapped Cache

16	10 000	Miss	000
3	00 011	Miss	011
16	10 000	Hit	000

Index	V	Tag	Data
000	Y	10	Mem[10000]
001	N		
010	Y	11	Mem[11010]
011	Y	00	Mem[00011]
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		



18	10 010	Miss	010
16	10 000	Hit	000

Index	V	Tag	Data
000	Y	10	Mem[10000]
001	N		
010	Y	10	Mem[10010]
011	Y	00	Mem[00011]
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

Associative Caches

- Fully associative
 - Allow a given block to go in any cache entry
 - Requires all entries to be searched at once
 - Comparator per entry (expensive)
- n-way set associative
 - Each set contains n entries
 - Block number determines which set
 - (Block number) modulo (#Sets in cache)
 - Search all entries in a given set at once
 - n comparators (less expensive)

Associative Cache Example


- Example: Placement of a block whose address is 12:

$$12 \bmod 8 = \text{block 4}$$




Direct mapped

$$12 \bmod 4 = \text{set 0, can be in either way of the set}$$



Set associative

can appear in any of the eight cache blocks



Fully associative

Spectrum of Associativity

- An eight-block cache configured as direct mapped, two-way set associative, and fully associative.

One-way set associative

(direct mapped)

Block	way,0 Tag	Data
0		
1		
2		
3		
4		
5		
6		
7		

Two-way set associative

Set	way,0 Tag	Data	way,1 Tag	Data
0				
1				
2				
3				

Four-way set associative

Set	way,0 Tag	Data	way,1 Tag	Data	way,2 Tag	Data	way,3 Tag	Data
0								
1								

Eight-way set associative (fully associative)

way 0		way 1		way 2		way 3		way 4		way 5		way 6		way 7	
Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data

For Internal Use Only!

Associativity Example

- 4-block caches, 1byte/block
 - Direct mapped, 2-way set associative, fully associative
 - Block access sequence: 0, 8, 0, 6, 8
 - Direct mapped
 - index = (Block address) mod (**#Blocks**)
- needs 2 bit for index

Block address	Cache index	Hit/miss	Cache content after access			
			way 0 block 0	way 0 block 1	way 1 block 2	way 1 block 3
0	0	miss	Mem[0]			
8	0	miss	Mem[8]			
0	0	miss	Mem[0]			
6	2	miss	Mem[0]		Mem[6]	
8	0	miss	Mem[8]		Mem[6]	

Associativity Example (cont.)

- 2-way set associative

Block access sequence: 0, 8, 0, 6, 8

- index = (Block address) mod (**#Sets**)

needs 1 bit for index

Block address	Cache index	Hit/miss	Cache content after access			
			way ₀	Set 0 way ₁	Set 1 way ₀	Set 1 way ₁
0	0	miss	Mem[0]			
8	0	miss	Mem[0]	Mem[8]		
0	0	hit	Mem[0]	Mem[8]		
6	0	miss	Mem[0]	Mem[6]		
8	0	miss	Mem[8]	Mem[6]		

- Fully associative (no index)

needs 0 bit for index

Block address		Hit/miss	Cache content after access			
			way ₀	way ₁	way ₂	way ₃
0		miss	Mem[0]			
8		miss	Mem[0]	Mem[8]		
0		hit	Mem[0]	Mem[8]		
6		miss	Mem[0]	Mem[8]	Mem[6]	
8		hit	Mem[0]	Mem[8]	Mem[6]	

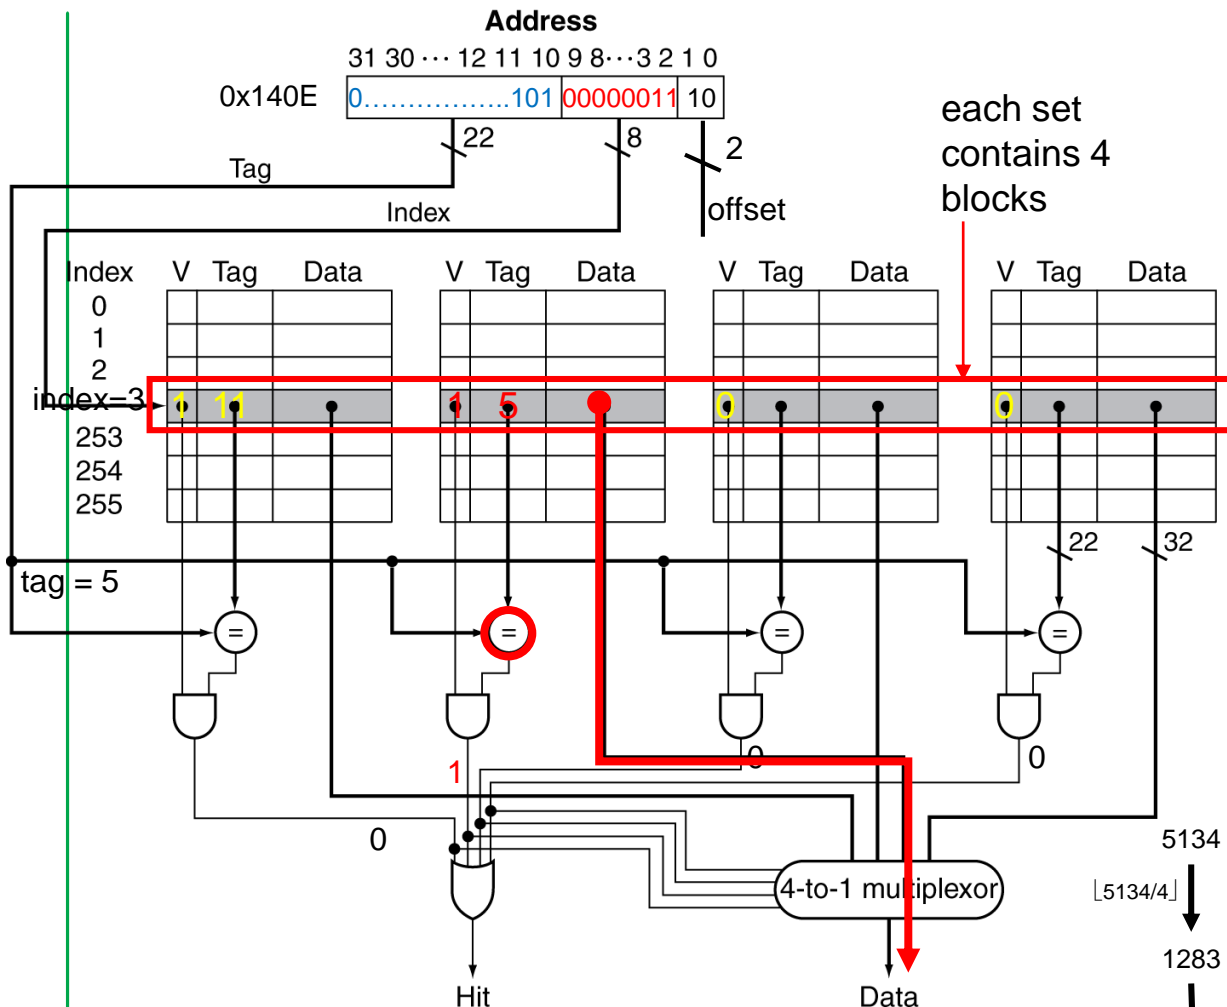
How Much Associativity

- Increased associativity decreases miss rate
 - But with diminishing returns
- Miss rate simulation of a system with 64KB D-cache, 16-word blocks, SPEC2000
 - 1-way: 10.3%
 - 2-way: 8.6%
 - 4-way: 8.3%
 - 8-way: 8.1%

Address Subdivision

- 4K blocks cache, 4-word/block, 32-bit address
- $\text{offset} = \log_2 16 = 4\text{bits} \rightarrow 28\text{bits for index} + \text{tag for all types}$
- direct mapped (1-way set associative)
 - same number of sets as blocks $\rightarrow 4096$ blocks(sets)
 - $\text{index} = \log_2 4096 = 12$ bits
 - $\text{tag} = 32 - 4 - 12 = 16$ bits
- 2-way set associative
 - 2 blocks/set $\rightarrow 2048$ sets
 - $\text{index} = \log_2 2048 = 11$ bits
 - $\text{tag} = 32 - 4 - 11 = 17$ bits
- 4-way set associative
 - 4 blocks/set $\rightarrow 1024$ sets
 - $\text{index} = \log_2 1024 = 10$ bits
 - $\text{tag} = 32 - 4 - 10 = 18$ bits
- fully associative
 - just 1 set \rightarrow no index
 - $\text{tag} = 32 - 4 = 28$

Set Associative Cache Organization



4KB cache, 1word/block
To what **set** number does address **0x140E** map?

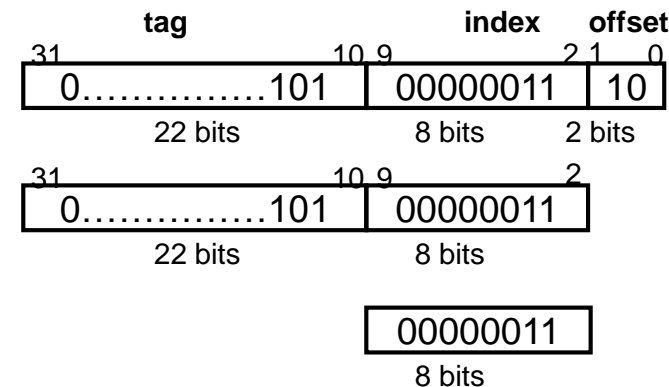
offset: $\log_2(4\text{byte}) = 2\text{bits}$

#blocks: $4\text{KB}/4\text{B} = 1\text{K blocks}$

#sets: $1\text{K blocks}/4 \text{ way} = 256 \text{ sets}$

index: $\log_2(256\text{sets}) = 8\text{bits}$

$0x140E_{\text{hex}} = \underbrace{101}_{\text{tag}} \underbrace{00000011}_{\text{index}} \underbrace{10}_{\text{offset}}$



Increasing associativity
shrinks index, expands tag



Replacement Policy


- Direct mapped: no choice
- Set associative
 - Prefer non-valid entry, if there is one
 - Otherwise, choose among entries in the set
- Least-recently used (LRU)
 - Choose the one unused for the longest time
 - Simple for 2-way, manageable for 4-way, too hard beyond that
- Random
 - Gives approximately the same performance as LRU for high associativity

Outline

- Measuring cache performance
- Improving performance – Associative cache
 - Fully associative
 - n-ways Set associative
- **Improving performance – Multilevel Caches**

Multilevel Caches

- Primary cache attached to CPU
 - Small, but fast
- Level-2 cache services misses from primary cache
 - Larger, slower, but still faster than main memory
- Main memory services L-2 cache misses
- Some high-end systems include L-3 cache



Multilevel Cache Example

- Given
 - CPU base CPI = 1, clock rate = 4GHz
 - Primary cache Miss rate/instruction = 2%
 - Main memory access time = 100ns
- Solution: With just primary cache
 - Miss penalty = $100\text{ns} / 0.25\text{ns} = 400$ cycles
 - Effective CPI = Base CPI + miss penalty/instruction
 $= 1 + 2\% \times 400 = 9$

Registers

L1 Cache
1 cycle

miss 2%

Memory
400 cycles

Example (cont.)

- Given: After adding L-2 cache
 - L2 cache Access time = 5ns
 - L2 global miss rate/instruction = 0.5%
- Solution 1, calculate based on **Global miss rate** (The fraction of references that miss in all levels):
 - L-1 miss (2%) first need to access the L2 cache
 - Penalty = $5\text{ns}/0.25\text{ns} = 20$ cycles
 - L-1 miss with L-2 also miss (0.5%)
 - Extra penalty = 400 cycles
 - Effective CPI = $1 + 2\% \times 20 + 0.5\% \times 400 = 3.4$
 - Speedup = $9/3.4 = 2.6$
- Solution 2, calculate based on **Local miss rate** (The fraction of references to one level of a cache that miss)
 - L2 local miss rate/instruction = $0.5\%/2\% = 25\%$
 - Effect CPI = $1 + 2\% \times (20 + 25\% \times 400) = 3.4$

Registers

L1 Cache
1 cycle

global miss 2%

L2 Cache
20 cycles

global miss 0.5%

Memory
400 cycles



Multilevel Cache Considerations

- Primary cache
 - Focus on minimal hit time
- L-2 cache
 - Focus on low miss rate to avoid main memory access
 - Hit time has less overall impact
- Results
 - L-1 cache usually smaller than a single cache
 - L-1 block size smaller than L-2 block size