



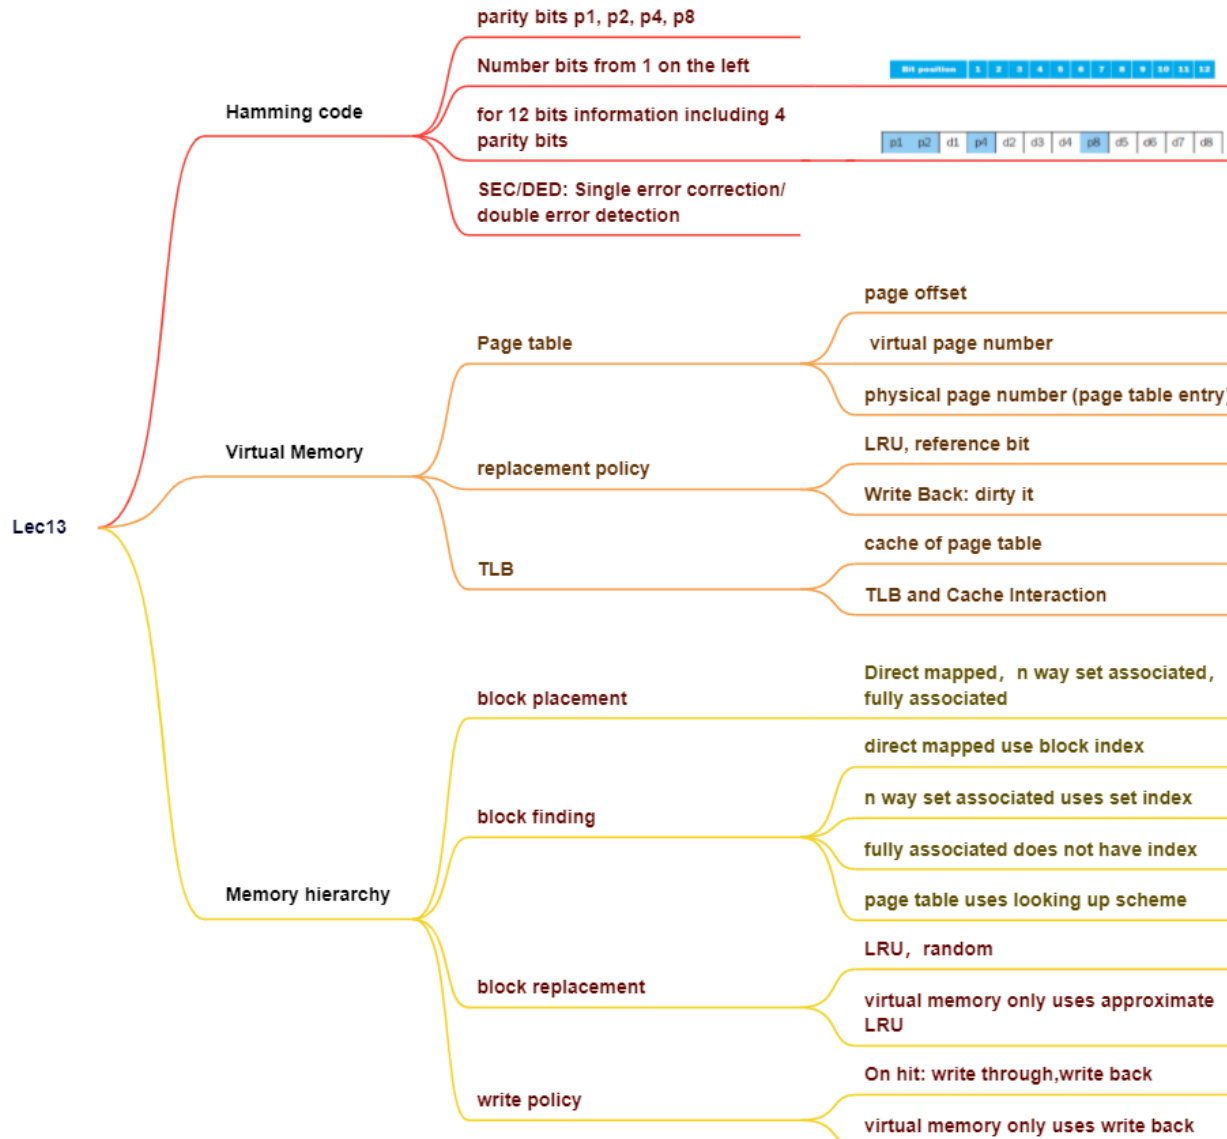
COMPUTER ORGANIZATION

Lecture 14 Parallel Processors from Client to Cloud

2024 Spring




Recap



Introduction

- Goal: replacing large inefficient processors with many smaller, efficient processors to get better performance per joule
 - Multiprocessors, cluster
 - Scalability, availability, power efficiency
- Task-level (process-level) parallelism
 - High throughput for independent jobs
- Parallel processing program
 - Single program run on multiple processors
- Multicore microprocessors
 - Chips with multiple processors (cores)
 - Shared Memory Processors (SMP)



Hardware and Software

- Challenge: hardware and software design that enables **parallel processing programs**, which can be **efficiently executed** (in performance and energy) when number of cores scales.
- Hardware
 - Serial: e.g., Pentium 4
 - Parallel: e.g., Core i7
- Software
 - Sequential: e.g., matrix multiplication
 - Concurrent: e.g., operating system
- Sequential/concurrent software can run on serial/parallel hardware

Parallel Programming

- We use “parallel processing program” to mean either sequential or concurrent software running on parallel hardware
- It’s hard to create parallel software
- Parallel programming needs to achieve significant performance improvement
 - Otherwise, just use a faster uniprocessor, since it’s easier!
- Difficulties of parallel programming:
 - Partitioning
 - Coordination
 - Communications overhead

Amdahl's Law

- Amdahl's Law: Calculate how much a computation can be sped up by running part of it in parallel

$$\text{Speedup} = \frac{T_1 + T_2}{T'} = \frac{T_1 + T_2}{\frac{T_1}{N} + T_2} = \frac{1}{\frac{1}{N} \left(\frac{T_1}{T_1 + T_2} \right) + \left(\frac{T_2}{T_1 + T_2} \right)}$$

$$\text{Speedup} = \frac{1}{\frac{1}{N}(F_{\text{parallelizable}}) + (1 - F_{\text{parallelizable}})}$$

$T_1 = T_{\text{affected}}$
 $T_2 = T_{\text{unaffected}}$
 $T' = T_{\text{improved}}$
 $N = \text{improve factor}$
 $F = \text{parallel portion}$

Amdahl's Law

- Sequential part can limit speedup
- Example: 100 processors, how to achieve 90× speedup?

- 1 processor: $F_{\text{sequential}} + F_{\text{parallelizable}} = 1 (T_{\text{old}})$

- 100 processors: $T_{\text{new}} = F_{\text{sequential}} + F_{\text{parallelizable}}/100$

$$\text{Speedup} = \frac{1}{(1 - F_{\text{parallelizable}}) + F_{\text{parallelizable}}/100} = 90$$

- $\rightarrow F_{\text{parallelizable}} = 0.999$
- sequential part should be no more than 0.1% of total task

Scaling Example

- Workload: sum of 10 scalars, sum of two 10×10 matrix
 - Assume the sum of 10 scalars cannot be paralleled
 - What is the speedup of 10 processors and 40 processors?
- Single processor: Time = $(10 + 100) \times t_{\text{add}}$
- 10 processors
 - Time = $10 \times t_{\text{add}} + 100/10 \times t_{\text{add}} = 20 \times t_{\text{add}}$
 - Speedup = $110/20 = 5.5$ ($5.5/10=55\%$ of potential)
- 40 processors
 - Time = $10 \times t_{\text{add}} + 100/40 \times t_{\text{add}} = 12.5 \times t_{\text{add}}$
 - Speedup = $110/12.5 = 8.8$ ($8.8/40=22\%$ of potential)
- Assumes load can be balanced across processors

Scaling Example (cont.)

- What if matrix size is 20×20 ?
- Single processor: Time = $(10 + 400) \times t_{\text{add}}$
- 10 processors
 - Time = $10 \times t_{\text{add}} + 400/10 \times t_{\text{add}} = 50 \times t_{\text{add}}$
 - Speedup = $410/50 = 8.2$ ($8.2/10=82\%$ of potential)
- 40 processors
 - Time = $10 \times t_{\text{add}} + 400/40 \times t_{\text{add}} = 20 \times t_{\text{add}}$
 - Speedup = $410/20 = 20.5$ ($20.5/40=51\%$ of potential)
- Assuming load balanced
- Make common case faster

Strong vs Weak Scaling

- Speedup:

task No. of cores	10 scalar 10*10 matrix	10 scalar 20*20 matrix
10	5.5 (55% of potential)	8.2 (82% of potential)
40	8.8 (22% of potential)	20.5 (51% of potential)

- Strong scaling: keep problem size fixed, time is reverse proportional to number of processors: $T(N,P)=T(1,P)/N$
- Weak scaling: constant time cost when problem size is proportional to number of processors:
 $T(N_1,P_1)=T(N_2,P_2)$
 - 10 processors (N_1), 10 × 10 matrix (P_1)
 - Time = $10 \times t_{\text{add}} + 100/10 \times t_{\text{add}} = 20 \times t_{\text{add}}$
 - 40 processors (N_2), 20 × 20 matrix (P_2)
 - Time = $10 \times t_{\text{add}} + 400/40 \times t_{\text{add}} = 20 \times t_{\text{add}}$
 - Constant performance in this example

Load Balancing

- In the above example
 - 40 processors are used to achieve 20.5 speedup
 - 40 processors are assumed to have balanced load (2.5% each)
- how about one processor with high load (5%)?
 - one processor takes $5\% \times 400 = 20$ adds, the others takes the rest $400 - 20 = 380$ adds
 - Time = $\max(380t/39, 20t/1) + 10t = 30t$
 - speedup = $410t/30t = 14$, smaller than 20.5
- how about one processor with higher load (12.5%)?

Parallel Processing

- The following techniques can enable parallel processing
 - SIMD, vector (section 6.3)
 - Multithreading (section 6.4)
 - SMPs and clusters (section 6.5)
 - GPUs (section 6.6)

SISD, MIMD, SPMD, SIMD and vector

		Data Streams	
		Single	Multiple
Instruction Streams	Single	SISD: Intel Pentium 4	SIMD: SSE of ISA x86
	Multiple	MISD: No examples today	MIMD: Intel Core i7

- SISD: single instruction stream, single data stream
 - Uniprocessor
- MIMD : multiple instruction, multiple data
 - Multi-core processor
 - SPMD: single program, multiple data
 - Typical way to write program on a multi-core processor
 - One program run on multiple processors
 - Different processors execute on different sections of code

SIMD

- Operate on vectors of data
 - Provide data level parallelism
 - E.g., MMX (MultiMedia eXtension) and SSE (Streaming SIMD Extension) instructions in x86
 - Multiple data elements in 128-bit wide registers
- All processors execute the same instruction at the same time
 - Each with different data address, etc.
- Simplifies synchronization
- Reduced instruction control hardware
- Works best for highly data-parallel applications

Vector Processors

- Highly pipelined function units
- Stream data from/to vector registers to units
 - Data collected from memory into registers
 - Results stored from registers to memory
- Example: Vector extension to RISC-V
 - v0 to v31: $32 \times$ 64-element registers, (64-bit elements)
 - Vector instructions
 - fld.v, fsd.v: load/store vector
 - fadd.d.v: add vectors of double
 - fadd.d.vs: add scalar to each element of vector of double
- Significantly reduces instruction-fetch bandwidth

An example using vector instruction

Calculate:

$$Y = a \times X + Y$$

➤ X and Y are vectors of 64 double precision numbers

➤ the starting addresses of X and Y are in \$s0 and \$s1

	l.d	\$f0,a(\$sp)	:load scalar a
	addiu	\$t0,\$s0,#512	:upper bound of what to load
loop:	l.d	\$f2,0(\$s0)	:load x(i)
	mul.d	\$f2,\$f2,\$f0	:a*x(i)
	l.d	\$f4,0(\$s1)	:load y(i)
	add.d	\$f4,\$f4,\$f2	:a*x(i) + y(i)
	s.d	\$f4,0(\$s1)	:store into y(i)
	addiu	\$s0,\$s0,#8	:increment index to x
	addiu	\$s1,\$s1,#8	:increment index to y
	subu	\$t1,\$t0,\$s0	:compute bound
	bne	\$t1,\$zero,loop	:check if done

Code using vector instructions:

	l.d	\$f0,a(\$sp)	:load scalar a
	lv	\$v1,0(\$s0)	:load vector x
	mulvs.d	\$v2,\$v1,\$f0	:vector-scalar multiply
	lv	\$v3,0(\$s1)	:load vector y
	addv.d	\$v4,\$v2,\$v3	:add y to product
	sv	\$v4,0(\$s1)	:store the result

Example: DAXPY ($Y = a \times X + Y$)

Conventional RISC-V code:

```

fld    f0,a(x3)      // load scalar a
addi   x5,x19,512    // end of array x
loop:  fld    f1,0(x19) // load x[i]
      fmul.d  f1,f1,f0  // a * x[i]
      fld    f2,0(x20) // load y[i]
      fadd.d  f2,f2,f1  // a * x[i] + y[i]
      fsd    f2,0(x20) // store y[i]
      addi   x19,x19,8  // increment index to x
      addi   x20,x20,8  // increment index to y
      bltu   x19,x5,loop // repeat if not done
  
```

Calculate:

$$Y = a \times X + Y$$

➤ X and Y are vectors of 64 double precision numbers

➤ the starting addresses of X and Y are in \$s0 and \$s1

Vector RISC-V code:

```

fld    f0,a(x3)      // load scalar a
fld.v  v0,0(x19)     // load vector x
fmul.d.vs v0,v0,f0    // vector-scalar multiply
fld.v  v1,0(x20)     // load vector y
fadd.d.v v1,v1,v0     // vector-vector add
fsd.v  v1,0(x20)     // store vector y
  
```



Multithreading

- Performing multiple threads of execution in parallel
 - Replicate registers, PC, etc.
 - Fast switching between threads
 - Threads: a lightweight process, share single address space
- Fine-grain multithreading
 - Switch threads after each cycle
 - Interleave instruction execution
 - If one thread stalls, others are executed
- Coarse-grain multithreading
 - Only switch on long stall (e.g., L2-cache miss)
 - Simplifies hardware, but doesn't hide short stalls (eg, data hazards)

Simultaneous Multithreading

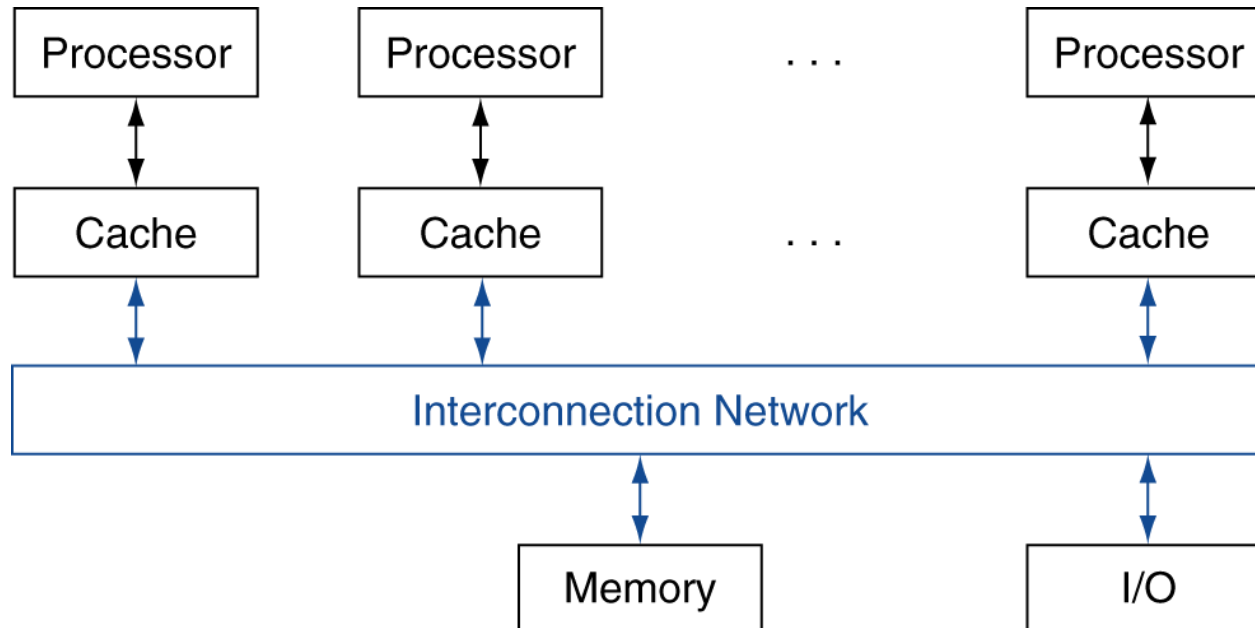
- In multiple-issue dynamically scheduled processor
 - Schedule instructions from multiple threads
 - Instructions from independent threads execute when function units are available
 - Within threads, dependencies handled by scheduling and register renaming
- Example: Intel Pentium-4 HT
 - Two threads: duplicated registers, shared function units and caches

Multithreading Example




Shared Memory Multiprocessors

- Multithreading improves efficiency of one processor
- How to efficiently programming on multiprocessor?
 - Don't want to rewrite old programs in single processor
 - Share memory among multiple cores
- SMP: shared memory multiprocessor



Message Parsing Multiprocessors

- Each processor has private physical address space
- Hardware sends/receives messages between processors



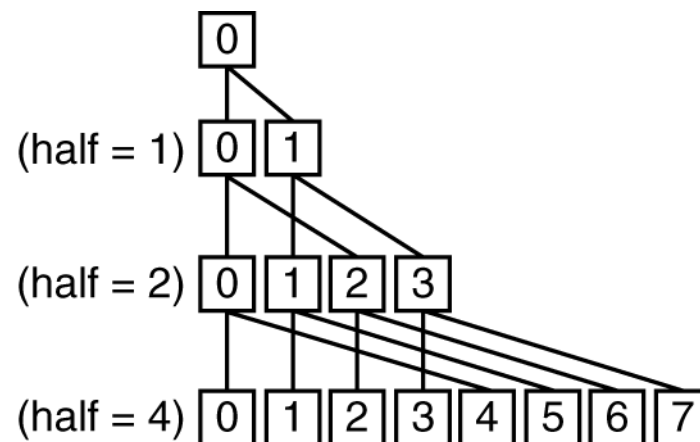
Loosely Coupled Clusters

- Network of independent computers
 - Each has private memory and OS
 - Connected using I/O system
 - E.g., Ethernet/switch, Internet
- Suitable for applications with independent tasks
 - Web servers, databases, simulations, ...
- High availability, scalable, affordable
- Problems
 - Administration cost (prefer virtual machines)
 - Low interconnect bandwidth
 - c.f. processor/memory bandwidth on an SMP

Example: Sum Reduction

- Sum 100,000 numbers on 100 processor
 - Partition 1000 numbers per processor
 - Initial summation on each processor

```
sum[Pn] = 0;  
for (i = 1000*Pn;  
     i < 1000*(Pn+1); i = i + 1)  
    sum[Pn] = sum[Pn] + A[i];
```
 - Now need to add these partial sums
 - Reduction: divide and conquer
 - Half the processors add pairs, then quarter, ...



Cloud Computing and Data Center

- Cloud computing
 - Warehouse Scale Computers (WSC)
 - Software as a Service (SaaS)
 - Portion of software run on a PMD (personal mobile device) and a portion run in the Cloud
 - Amazon and Google
- Data centers
 - Millions of computers connected by off-the-shelf networking devices

Google Data Centers



Introduction of GPUs

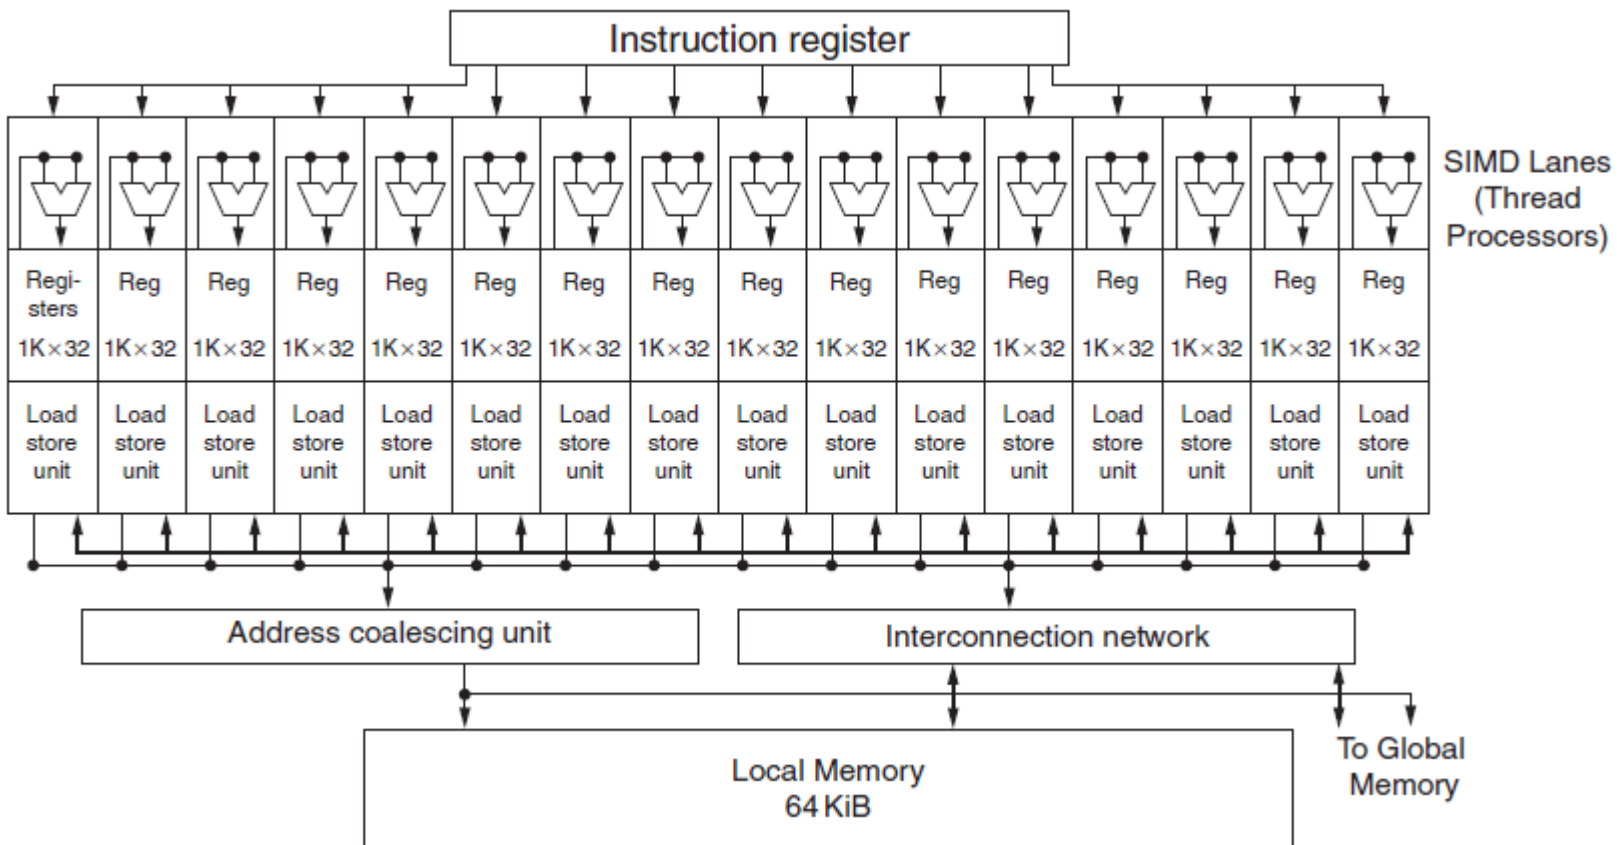
- Development of GPU
 - General-purpose CPU not suitable for graphic processing
 - Game industry drive the improvement of graphic processing
 - GPU (graphic processing unit) appear and developed faster than general-purpose CPU, high performance, low cost
 - Easy-use programming language helps GPU's popularity
- Difference between GPU and CPU
 - GPU supplement CPU, doesn't replace it. GPU doesn't need to perform all tasks of CPU
 - GPU problem size is MB to GB, not hundreds of GB.

GPU Architectures

- Processing is highly data-parallel
 - GPUs are highly multithreaded, 16-32 threads
- GPU memory oriented towards bandwidth rather than latency
 - Graphics memory is wide and high-bandwidth
 - GPU memory is smaller than CPU. 4-6GB, instead of 256GB or above
 - Less reliance on multi-level caches
- Trend toward general purpose GPUs
 - Heterogeneous CPU/GPU systems
 - CPU for sequential code, GPU for parallel code
- Programming languages/APIs
 - DirectX, OpenGL
 - C for Graphics (Cg), High Level Shader Language (HLSL)
 - Compute Unified Device Architecture (CUDA)

Block Diagram of a SIMD processor in GPU

- A GPU consists of multiple multi-thread SIMD processors (Nvidia Fermi architecture)



Classifying GPUs

	Static: Discovered at Compile Time	Dynamic: Discovered at Runtime
Instruction-Level Parallelism	VLIW	Superscalar
Data-Level Parallelism	SIMD or Vector	Nvidia Multiprocessor

xPUs

- CPU: good for control, sequential programming
- GPU: good for graphics, parallel programming
 - CPU/GPU mixed architecture
- TPU: Tensor processing unit
 - Proposed by Google, targeting at acceleration for tensorflow platform
Suitable for machine learning model training and testing
- DPU: deep learning processing unit
 - Proposed by DeePhi Tech, FPGA-based processing unit
- NPU: neural network processing unit
 - IBM TrueNorth
- BPU: brain processing unit

Concluding Remarks

- Goal: higher performance by using multiple processors
- Difficulties
 - Developing parallel software
 - Devising appropriate architectures
- From multicore to data centre
- Performance per dollar and performance per Joule drive both mobile and WSC



Summary for this course

- Digital logic → computer organization → operating system/embedded system
- Main content:
 - Processor
 - Memory
 - Parallel
- Hardware thinking is important: resource tradeoff